

Concurrent Threads (1)

Recommended Reading:

Bacon, J.:	Concurrent Systems (9, 10, 11)
Silberschatz, A.:	Operating System Concepts (4, 6, 7)
Stallings, W.:	Operating Systems (5, 6)
Tanenbaum, A.:	Modern Operating Systems (2)
Wettstein, H.:	Systemarchitektur (7, 8, 9, 10)

Roadmap for the next lectures

Problems with Concurrency

Interrelationship Patterns

Well known Concurrency Problems

Signaling and Synchronization

Mutual Exclusion

Communication

Deadlocks

Transactions

Potential problems with Concurrency

Concurrent threads often need to **share resources** and/or **data** (maintained either in shared memory or files)

- If there is no controlled access to shared data, threads may get **inconsistent** data
- The overall result of an application performed by concurrent threads **may depend** on the execution sequence of these threads, i.e. a **race condition**
- Concurrent threads may even induce **conflicts** due to competition around exclusive resources

Simple Control Problem

Assumption: Given a shared resource (suppose a display)

```
{Thread 1}
boolean c1:=true
while c1 do

    for i=1 step 1 until MAXINT
        display("i ")

    {do something else}
    begin ... c1:=false ... end
od
```

Expected Behavior: Until c1= false thread 1 displays sequences of
1 2 3 4 ... MAXINT onto the screen

Simple Control Problem

What may happen now?

Assumption: Given a shared resource (suppose a display)

```
{Thread 1}
boolean c1:=true
while c1 do

    for i=1 step 1 until MAXINT
        display("i ")

    {do something else}
    begin ... c1:=false ... end
od
```

```
{Thread 2}
boolean c2:=true
while c2 do

    for i=1 step 1 until MAXINT
        display("-i ")

    {do something else}
    begin ... c2:=false ... end
od
```

Output Mix:

1 2 3 4 5 6 7 8 -1 -2 -3 -4 -5 -6 -7 -8 9 10 ...

due to various dispatching activities (e.g. end of time slice)

Problem: Serialization of Critical Sections

Assumption: Given a shared resource (suppose a display)

```
{Thread 1}  
boolean c1:=true  
while c1 do
```

```
  for i=1 step 1 until MAXINT  
    display("i ")
```

```
  {do something else}  
  begin ... c1:=false ... end  
od
```

```
{Thread 2}  
boolean c2:=true  
while c2 do
```

```
  for i=1 step 1 until MAXINT  
    display("-i ")
```

```
  {do something}  
  begin ... c2:=false ... end  
od
```

Red Boxes = Critical Program Sections because both threads access the **shared display**

Consequence: We have to offer something enabling **serialization** of critical sections

Another Concurrency Problem

```
integer a,b :=1;  // shared data for both threads
```

```
{Thread 1}
```

```
while true do
```

```
  a = a + 1;  
  b = b + 1;
```

```
  {do something else}  
od
```

```
{Thread 2}
```

```
while true do
```

```
  b = b + 2;  
  a = a + 2;
```

```
  {do something else}  
od
```

Both Threads read (and write to) the shared global data a, b
=> data inconsistency, **a != b** after some time due to dispatching!

Conclusion: Resource and data sharing may lead to similar problems!

Interrelationship Patterns with Concurrent Threads


Degree of Awareness	Relationship	Mutual Influence	Control Problems
Unaware of each other	Competition	<ul style="list-style-type: none">• Results are independent• Timing may be affected	<ul style="list-style-type: none">• Starvation
Indirectly aware of each other (shared object)	Cooperation by Sharing	<ul style="list-style-type: none">• Results may be dependent• Timing may be affected	<ul style="list-style-type: none">• Mutual Exclusion• Deadlock• Starvation• Data Coherence
Directly aware of each other (communication)	Cooperation by communication	<ul style="list-style-type: none">• Results may be dependent• Timing may be affected	<ul style="list-style-type: none">• Deadlock• Starvation

Some Classical Concurrency Problems

- Producer/Consumer Problem
- Barbershop Problem
- Reader/Writer Problem
- Cigarette Smoker Problem
- Monkey Rock Problem
- Dining Philosophers

Remark: Discuss them **all** very carefully and find out different solutions, i.e. solutions based upon different **synchronization mechanisms** (see later)!


Producer/Consumer Problem with unbounded buffers



Do we really have a concurrency problem with an unbounded buffer, if there are only 1 producer and 1 consumer?

Of course, a consumer cannot consume from an empty buffer.


Producer/Consumer Problem with bounded buffers



Additional problems with a bounded buffer?

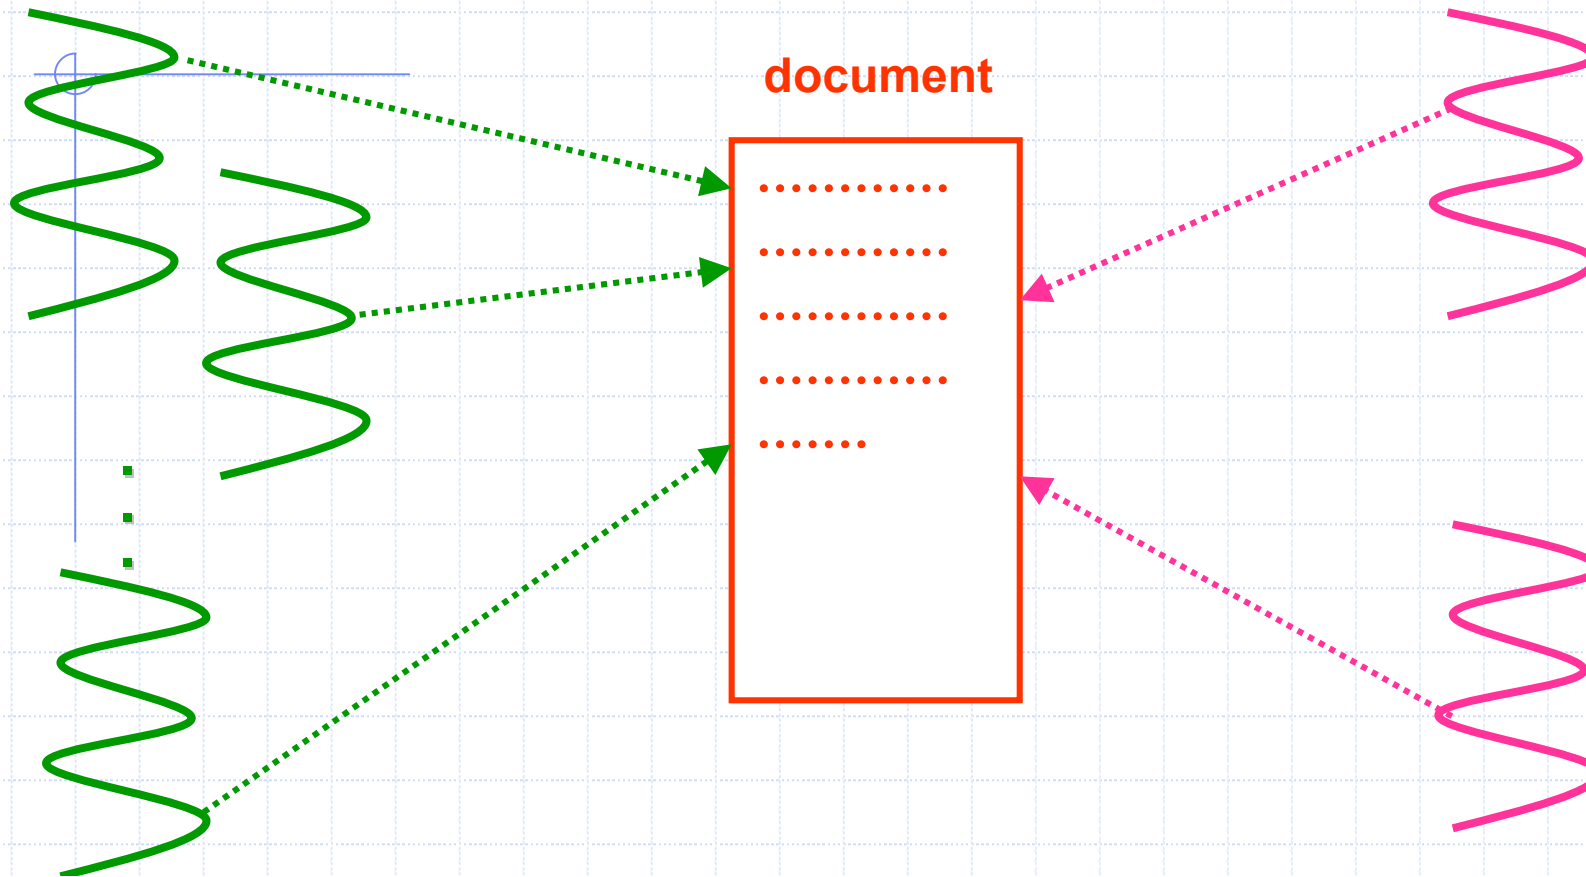
Additional problems with $p > 1$ producer and/or respectively $c > 1$ consumer?

Barber Shop Problem



Concurrent Threads

Reader/Writer Problem



Which Problems may occur?

Data Consistence


Patil's Smoker Problem

Given: A vendor store for smoking ingredients
3 chain smokers and the sleeping vendor
For smoking a smoker needs: tobacco, paper and matches
Smoker A has his own tobacco
Smoker B has his own paper
Smoker C has his own matches


Each time the sleeping vendor is awakened he puts two smoking ingredients on his table allowing one of the smokers to continue his unhealthy pleasure. After smoking, the smoker wakes up the vendor again who then puts another pair of smoking ingredients (at random) thus unblocking another smoker.

Monkey Rock Problem

The northern and southern monkeys have to eat and drink!
There is a small rope between the two rocks,
but there are also hungry crocodiles below the rope.
The rope can carry up to $m \geq 1$ monkeys, concurrent
crossing in both direction is not possible.



Dijkstra's famous Dining Philosopher Problem



Life of a philosopher:


```
repeat forever
begin
  thinking
  getting hungry
  getting 2 forks
  eating
end
```

Requirements:

No Deadlock!

No Starvation!

Dijkstra's Dining Philosophers - solution I



Life of a philosopher:


```
do
  think
  get hungry
  acquire 2 forks
  eat
  release forks
od
```

```
think: do
  if neighbor requests fork
    then give it to him
  fi
until hungry od.
```

```
hungry: do
  if fork missing
    then request it from neighbor
  fi
  if neighbor requests fork
    then give it to him
  fi
until have both forks od.
```

deadlock free?
starvation

Dijkstra's Dining Philosophers - solution 1a



Life of a philosopher:

```
do
  think
  get hungry
  acquire 2 forks
  eat
  release forks
od
```


```
think: do
  if neighbor requests fork
    then give it to him
  fi
until hungry od
```

```
hungry: do
  if fork missing
    then request it from neighbor
    if got dirty fork
      then clean it
    fi
  fi
  if neighbor requests dirty fork
    then give it to him
  fi
until have both forks od.
```

deadlock free ?

no starvation ?


Dijkstra's Dining Philosophers - solution Ia



conclusion from solution I:

- need communication operations for communicating sequential processes, e.g.
 - send request
 - receive request

Dijkstra's Dining Philosophers - solution II



Life of a philosopher:

```
do
  think
  get hungry
  acquire 2 forks
  eat
  release forks
od
```


atomic operation

deadlock free

starvation

```
Enter critical section
get left fork
get right fork
leave critical section
```


Dijkstra's Dining Philosophers - solution II



conclusion from solution II:

- need synchronization operations on shared data objects, e.g.

?



Signaling and Synchronization

Long History of Signaling Techniques

- **Drums** (Bush Men)
- **Smoke Signals** (Red Indians)
- **Signal Fires** (e.g. in War Times)
- **Signal Flags** (e.g. Marine)
- **Signal Whistles** (e.g. Sports)
- **Signal Colors** (e.g. Traffic Lights, “Mimikri”)
- **Signal Perfumes** (e.g. Mammals)
- ...

Potential meanings of a Signal

- “Pay Attention” (see a siren)
- “Stop” (see road signs)
- “Go Ahead” (see whistle of a station officer)
- “Interrupt” (see whistle of an arbiter in a soccer game)
- ...

=> You must know the exact meaning of each signal
otherwise your reaction on a signal may be wrong

Implementing a Signal

- **Flag**

1 = Signal set, 0 = Signal reset

- **Continuation** (signaled thread may continue)
- **Stop** (signaled thread has to wait)
- **Abort** (signaled thread has to be aborted)

...

see “signal vector” in Unix or Linux

- **Counter**

Any value may have a different meaning or just reflects the number of pending signals

Problem: Try to find out when a flag is sufficient and when you'll need a counter!

Synchronizing a Precedence Relation

{Thread 1}

•

•

•

{section a1
... }

{ section a2
... }

•

•

•

{Thread 2}

•

•

•

{ section b1
... }

{ section b2
... }

•

•

•

Problem: *How to achieve that $a1 <^* b2$ ($a1$ precedes $b2$), i.e. section $b2$ cannot be executed before section $a1$ has completed?*

Implementing a Precedence Relation

```
1:1_signal s;
```

```
{Thread 1}
```

```
•
```

```
•
```

```
•
```

```
{section a1
```

```
... }
```

```
Signal(s)
```

```
{ section a2
```

```
... }
```

```
•
```

```
•
```

```
•
```

```
{Thread 2}
```

```
•
```

```
•
```

```
•
```

```
{ section b1
```

```
... }
```

```
Wait(s)
```

```
{ section b2
```

```
... }
```

```
•
```

```
•
```

```
•
```

Problem: *How to implement a 1:1_signal_object?*

Principal Types of Solutions

Software solutions

algorithms neither relying on special hardware nor OS features

Hardware solutions

relying on some special machine instructions

Operating System solutions

providing kernel functions to system and application programmers

Remark: Most systems offer just a subset of the above types of solutions.

Software Solutions

Flag s as common variable of both threads

signal(s)

set s



wait(s)

s set ?

no

reset s



What happens if 'signal' is invoked prior to 'wait'?

What may happen if 'wait' is invoked prior to 'signal'?

Discuss this "proposal" carefully!
Does it work on any system?
Is it effective and/or efficient?

Implementation of a 1:1 signal in software, version I1

```
module 1:1_signal
  export Signal, Wait;
  import yield;
  type signal = record
    S: signal := reset
  end
```

```
  procedure Signal(SO:signal)
  begin
    SO.S := set;
    yield(); /* anonymous yield */
  end
```

```
  procedure Wait (SO:signal)
  begin
    while SO.S = reset do
      yield()
    od
    SO.S :=reset
  end
end module
```

Requires cooperative scheduling

No signal loss ?
Efficient ?
Scalable ?

Implementation of a 1:1 signal, version II

```
module 1:1_signal
  export Signal, Wait
  import block, unblock
  type signal = record
    W: thread := nil
  end

  procedure Signal(SO:signal)
  begin
    unblock (SO.W);
  end

  procedure Wait (SO:signal)
  begin
    SO.W := myself ;
    block (myself) ;
    SO.W := nil
  end
end module
```

No signal loss?

Efficient !

Scalable !

Synchronization

Implementing a 1:1 signal, version III

```
module 1:1_signal
  export Signal, Wait
  import UnblockThread, BlockThread
  type signal = record
    S: signal := reset
    W: waiting thread := nil
  end

  procedure Signal(SO:signal)
  begin
    SO.S := set;
    if SO.W ≠ nil then
      UnblockThread(SO.W)
    end

  procedure Wait (SO:signal)
  begin
    while SO.S = reset do
      SO.W := myself
      BlockThread(myself)
    od
    SO.S := reset
  end

end module
```

No signal loss !
Efficient !

Scalable !

Race condition !

Open Problem

Conclusion:

Operations wait and signal should be atomic!
(We'll postpone how to achieve this property.)

Synchronization of a Mutual Precedence Relation

{Thread 1}

•
•
•

{section a1
... }

{section a2
... }

•
•
•

{Thread 2}

•
•
•

{section b1
... }

{ section b2
... }

•
•
•

Problem: *How to achieve $a1 <^* b2$ and $b1 <^* a2$?*

Solution of a Mathematician

Idea: Transform the new problem to an already solved one and take the known solution

```
1:1_signal s1,s2;
```

```
{Thread 1}
```

```
·
```

```
·
```

```
·
```

```
{section a1
```

```
... }
```

```
Signal(s1)
```

```
Wait(s2)
```

```
{section a2
```

```
... }
```

```
·
```

```
·
```

```
·
```

```
{Thread 2}
```

```
·
```

```
·
```

```
·
```

```
·
```

```
{section b1
```

```
... }
```

```
Signal(s2)
```

```
Wait(s1)
```

```
{ section b2
```

```
... }
```

```
·
```

```
·
```

```
·
```

Remark: Discuss Pros and Cons!

Mutual Precedence Relations via signal and wait

Pros:

No extra Mechanism for a related Problem

Cons:

Complicated for $n > 2$ Threads

Low Performance due to many Kernel Calls

Let's invent a new mechanism with a better behavior!

Synchronize Operation for Mutual Precedence Relations

`sync s`

`{Thread 1}`

`.`

`.`

`.`

`{section a1`
`... }`

Synchronize(s)

`{section a2`
`... }`

`.`

`.`

`.`

`{Thread 2}`

`.`

`.`

`.`

`{section b1`
`... }`

Synchronize(s)

`{section b2`
`... }`

`.`

`.`

`.`



Problem: *How to implement a synchronization module for two threads ?*

Synchronization

Synchronize Operation for Mutual Precedence Relations

```
module synchronization
  export synchronize
  import UnblockThread, BlockThread
  type sync = record
    S: signal := reset
    W: waiting thread := empty
  end

  procedure Synchronize(SY:sync)
    begin
      if SY.S = reset
      then begin
        SY.S := set
        SY.W := myself
        BlockThread(myself)
      end
      else begin
        UnblockThread(SY.W)
        SY.S := reset
      end
    end
  end module
```

{I am first}

{and wait for my partner}

{I am second and}

{release my partner and}


{do a reset for future reuse}

Exercise: Generalize this module for $n > 2$ threads!

Application of a Multiple Synchronize Operation


```
. {some numerical problem solved via difference equations}
.
.
while true do
  begin
    for all i,j
      begin
        temp[i,j] := old[i-1,j] + old[i+1,j]
      end
    n_synchronize(S)
    for all i,j
      begin
        old[i,j] := temp[i,j]
      end
    n_synchronize(S)
  end
.
.
.
```

More Patterns on Signaling: Many to One




The pink thread may continue iff the 3 other threads passed certain code sections.

Example for a Many to One Signaling Relation



These 4 threads form a team working together on some application problem.
3 production threads calculate intermediate results (buffering in shared memory).
Any of the productions threads signals when he finished another calculating step.
The display thread, averaging these 3 intermediate results, displays the average.


More Patterns on Signaling: One to Many



The 2 pink threads may continue iff the green thread passed a certain code section

Think over an application for this pattern!


More Patterns on Signaling: Many to Many



2 pink threads may continue iff 3 thread passed certain code sections

Think over an application for this pattern!

Further Patterns of Signaling



It's sufficient for the pink thread that one of the 2 signaling green threads has delivered a signal.

Additional Problem: How to buffer signals?

Buffering Signals

- Each incoming signal is buffered until a potential waiting thread consumes this signal

Pro: Reaction on each signal

Con: A deficient signaling source may flood the system

- An incoming signal may overwrite a previous one (e.g. binary semaphore)

Pro: Reaction only on the newest signal

Con: Danger of lost signals

Dijkstras (Counting) Semaphores

A semaphore S is an integer variable that, apart from initialization, can only be accessed by 2 **atomic and mutually-exclusive** operations:

$P(S)$ $P \sim$ Passeren (some dutch signaling language)

$V(S)$ $V \sim$ Verlaaten (see above)

To avoid **busy waiting**: when a thread cannot “passeren” inside of $P(S)$ it will be put into a **blocked queue** waiting for an event to be done inside of $V(S)$ by some other thread (hopefully, otherwise it would starve).

Dijkstras (Counting) Semaphores

Semantic of a Counting Semaphore for signaling purpose:

1. A **positive** value of the counter indicates how many signals are pending
2. A **negative** value of the counter indicates how many threads are waiting for a signal, i.e. are queued within the semaphore object
3. If the **counter = 0** => no thread is waiting for a signal

Dijkstra's: Counting Semaphores (historical 1. solution)

```
module semaphore
  export P,V
  import BlockThread, UnblockThread
  type semaphore = record
    Count: integer := 0           {no signal is pending}
    QWT: list of Threads = empty {queue of waiting threads for sema}
  end

  P(S:semaphore):
    S.Count := S.Count - 1
    if S.Count < 0 then
      insert (S.QWT, myself)      {1 more waiting thread}
      BlockThread(myself)
    fi .

  V(S:semaphore)
    S.Count := S.Count + 1        {1 additional signal}
    if S.Count <= 0 then
      unblockThread(delete first(S.QWT))
    fi .

end
```


Mutual Exclusion

Dijkstra's: Counting Semaphores

Remark: Semaphores can also solve the next class of coordination problems:

Critical Sections

The Critical Section Problem (1)

When a thread executes code that manipulates shared data (or resource), we say that the thread is in its **critical section** (CS) (for that shared data or resource)

The execution of critical sections must be **mutually exclusive**: at any time, only one thread is allowed to execute its critical section (even with multiple CPUs)

Then each thread must request the permission to enter its critical section (CS)

The Critical Section Problem (2)

The section of code implementing this request is called the **entry section**

The **critical section** (CS) should be closed by a **leave section**

The remaining code (outside of CS) is the **remainder section** (RS)

What do we really need to solve critical section problems?

We have to establish and offer a **serialization protocol** **so** that the results of the threads will not depend on the order in which their execution has been interleaved.

Framework for the Analysis of the Solutions

Each thread executes at nonzero speed but no assumption on the relative speed of n threads
General structure of a thread:

Multiple CPUs may be present but memory hardware prevents simultaneous access to the same memory location.
No ordering assumptions for interleaved executions

```
{A general pattern of all threads participating  
in the critical section problem}  
repeat  
    enter section  
    CS {critical section}  
    exit section  
    RS {remainder section}  
forever
```

3 Requirements for a Valid Solution

Mutual Exclusion:

At any time, **at most one** thread can be in its critical section (CS)

Progress

If no thread is executing in its CS while some threads wishes to enter, only threads that are **not** in their RS can participate in the decision: which thread will be the next. This selection cannot be postponed indefinitely

Bounded Waiting

After a thread has made a request to enter it's CS, there is a bound on the number of times that the other threads are allowed to enter their CS otherwise the thread will suffer from **starvation**

Different Types of Solutions

Software solutions

algorithms neither relying on special hardware nor OS features

Hardware solutions

relying on some special machine instructions

Operation System solutions

provide some kernel functions to the programmers

Software solutions

First consider only 2 threads

Algorithm 1 and 2 will be uncomfortable respectively incorrect

Algorithm 3 is correct (Peterson's algorithm)

Then consider a solution for n threads

Bakery algorithm

Notation

We start with 2 threads: T_0 and T_1

When presenting thread T_i , T_j always denotes the other thread ($i \neq j$)

Algorithm 1

The shared variable **turn** is initialized (0 or 1) before executing any T_i
 T_i 's critical section is executed iff $\text{turn} = i$
 T_i is **busy waiting** if T_j is in CS \Rightarrow mutual exclusion is satisfied
Progress requirement is not satisfied since it requires strict alternation of CSs

```
thread  $T_i$ :  
repeat  
    while ( $\text{turn} \neq i$ ) { } ;  
        CS $i$   
         $\text{turn} := j$  ;  
        RS $i$   
forever
```

Analysis: Suppose T_0 has a large RS_0 , whereas T_1 has a small RS_1 .
If $\text{turn}=0$, T_0 may enter CS_0 , leaves it ($\text{turn}=1$), then executes its long RS_0 .

Meanwhile T_1 was also in CS_1 , leaves it ($\text{turn}=0$), then executes its RS_1
Then T_1 tries in vain to enter CS_1 ! T_1 must **wait** until T_0 leaves its long RS_0

Additional Requirement: The length of an RS shouldn't have any influence on the execution sequence of the participating threads!

Algorithm 2

Keep 1 Boolean variable for each thread: $\text{flag}[0]$ and $\text{flag}[1]$
 T_i signals that it is ready to enter its CS by setting: $\text{flag}[i] := \text{true}$
Mutual Exclusion is satisfied but not the progress requirement

```
thread  $T_i$ :  
repeat  
     $\text{flag}[i] := \text{true};$   
    while ( $\text{flag}[j]$ ) {};  
    CS  
     $\text{flag}[i] := \text{false};$   
    RS  
forever
```

Analysis: Suppose the following execution sequence holds:
 $T_0: \text{flag}[0] := \text{true}$
 $T_1: \text{flag}[1] := \text{true}$

What will happen?

Both threads will **wait forever**, none of them can enter its CS (**“DEADLOCK”**)

Algorithm 3 (Peterson's solution)

Initialization: $\text{flag}[0] := \text{flag}[1] := \text{false}$,
and $\text{turn} := 0$ (or 1)
Willingness to enter CS specified by
 $\text{flag}[i] := \text{true}$
If both threads attempt to enter their
CS simultaneously, only one turn
value will last
Exit section: specifies that T_i is
unwilling to enter CS

```
thread  $T_i$ :  
repeat  
     $\text{flag}[i] := \text{true};$   
     $\text{turn} := j;$   
do {} while  
    ( $\text{flag}[j]$  and  $\text{turn} = j$ ) ;  
    CS  
     $\text{flag}[i] := \text{false};$   
    RS  
forever
```

“Proof” of correctness for Algorithm 3

- To prove that mutual exclusion is preserved:

- T_0 and T_1 are both in their CS only if $flag[0] = flag[1] = true$ and only if $turn = i$ for each T_i (which is impossible)

- To prove that the progress and bounded waiting requirements are satisfied:

- T_i cannot enter CS only if stuck in ‘*while()*..’ with condition ‘*flag[j]*’ and ‘*turn = j*’.

- If T_j is not ready to enter CS then ‘*! flag[j]*’ and T_i can enter its CS

- If T_j has set ‘*flag[j]*’ and is in its ‘*while()*..’, then either *turn=i* or *turn=j*

- If *turn=i*, then T_i enters CS. If *turn=j* then T_j enters CS, but it will reset *flag[j]* on exit: allowing T_i to enter CS

- but if T_j has time to set *flag[j]*, it must also set *turn=i*

- since T_i does not change value of *turn* while stuck in ‘*while()*..’, T_i will enter CS after at most one CS entry by T_j (bounded waiting)

What about Faulty Threads?

If all 3 main criteria (mutual exclusion, progress, bounded waiting) are satisfied, then a **valid solution** will provide **robustness against failures** within the remainder section (RS) of a thread

Failures within RS are like infinitely long RS, i.e. they should not affect the other thread(s)

However, no **valid solution** can ever provide any robustness, if a thread fails within its critical section (CS). *Why?*

A Thread failing within its critical section may never perform **exit_section**, i.e. neither a signal-operation nor a V(S) nor something comparable thus affecting the other thread(s) severely!!

Bakery Algorithm for n threads

Before entering their CS, each T_i receives a number. The holder of the smallest number enters its CS (as it is at least in English bakeries, ...)

If T_i and T_j receive the same number:

if $i < j$ then T_i is served first, else T_j is served first

T_i resets its number to 0 in its exit section

Notation:

$(a,b) < (c,d)$ if $a < c$ or if $a = c$ and $b < d$

$\max(a_0, \dots, a_k)$ is a number b such that: $b \geq a_i$ for $i=0, \dots, k$

Shared data:

choosing: array[0.. $n-1$] of boolean; initialized to false

number: array[0.. $n-1$] of integer; initialized to 0

Correctness relies on the following fact:

If T_i is in CS and T_k has already chosen its number[k] $\neq 0$,
then $(\text{number}[i], i) < (\text{number}[k], k)$

Bakery Algorithm

```
thread Pi:
repeat
    choosing[i] := true;
    number[i] := max(number[0] .. number[n-1]) + 1;
    choosing[i] := false;
    for j := 0 to n-1 do {
        while (choosing[j]) {};
        while (number[j] != 0
            and (number[j], j) < (number[i], i)) {};
    }
    CS
    number[i] := 0;
    RS
forever
```

Drawbacks of Software Solutions

Threads requesting to enter their critical sections are **busy waiting** (consuming processor time needlessly)

- If the critical sections are long enough, it would be more efficient to **block** threads.

Hardware solutions: *interrupt disabling*

Single processor: mutual exclusion is preserved but efficiency of execution is degraded: while in CS, we cannot interleave execution with other threads being in their RS

- Multi processor: **mutual exclusion** is **not preserved at all**
- Delay of interrupt handling may affect the whole System
- Application programmers may abuse this facility

```
thread Ti:
repeat
    disable interrupts
    critical section
    enable interrupts
    remainder section
forever
```

Summary: In general not an acceptable solution!

Hardware solutions: *special instructions*

(1) some machines offer instructions that perform *read-modify-write* operations *atomically* (indivisible, on the same memory location):

inc	[mem]	
xchg	[mem], reg	
bts	[mem]	{bit test and set}

- (2) some machines offer *conditional LD/ST* instructions instead:
 - LDL [mem] processor becomes sensitive for memory address *mem*
 - STC [mem] fails if another processor executed STC on the same address in the meantime
- instructions like (1) execute mutually exclusive on multiple CPUs
- like (2) permit to emulate mutually-exclusive instructions

The test-and-set instruction

An algorithm that uses testset for mutual exclusion:
Shared variable b is initialized to 0
Only the first T_i -having set b - can enter its CS

```
boolean testset(int& i)
{ if (i=0) {
    i:=1;
    return true;
  } else
    return false;
}
```

```
thread  $T_i$ :
repeat
  repeat{ }
  until testset( $b$ ) ;
  CS
   $b:=0$  ;
  RS
forever
```

Analysis of the test-and-set solution (1)

Mutual exclusion is preserved: if T_i enters its CS, the other T_j perform **busy waiting**. Hence there is an **efficiency problem**

When T_i exits its CS, the selection of the T_j who will enter CS is arbitrary: **no bounded waiting**. Hence **starvation** is possible!

Processors (ex: Pentium) often provide an atomic `xchg(a,b)` or *compare-and-swap* instruction that swaps the content of `a` and `b`.

(But `xchg(a,b)` suffers from the same drawbacks as test-and-set)

*However, what's the biggest problem with one of these “**spin lock**” solutions?*

Analysis of the test-and-set solution (2)

Repeated test-and-set-instructions may monopolize the system bus affecting each other system activity (if related to that CS or if not)

This is a severe danger for another sort of “**deadlock**” on a Single Processor System

Solutions with OS/PL Support

- **Dijkstras (Counting) Semaphores**
- **Dijsktras (Binary) Semaphores**
- **Monitors**
- **Transactions**

Dijkstras (Counting) Semaphores

Semantic of a Counting Semaphore for **mutual exclusion** of critical sections:

1. A **positive** value of the counter indicates how many threads still may enter the critical section
2. A **negative** value of the counter indicates how many threads are waiting in front of the critical section, i.e. are queued within the semaphore object
3. If the **counter = 0** => no thread is waiting respectively maximally allowed threads are in the critical section

Still the open problem:

How can we get “atomic semaphore-operations” ?

Atomic Semaphore Operations

Problem:

$P()$ and $V()$ each consisting of multiple machine instructions have to be atomic!

Solution:

Need “another” **sort of critical sections**, hopefully with shorter execution times, to establish atomic semaphore operations!!

“very short”
enter_section

$P(S)$

“very short”
leave_section

Revisiting Dijkstras (Counting) Semaphores

The critical sections required for implementing $P(S)$ and $V(S)$ are very short: typically only 10 instructions.

Solutions (to establish these short critical sections around $P(S)$ or $V(S)$):

- single processor: disable interrupts during those operations

- multi processor: use previous software or hardware schemes

Atomic Counting Semaphores (Single Processor)

```
P(sema S)
begin
  DisableInterrupt
  s.count--
  if s.count < 0 then
    BlockThread(S)
  fi
  EnableInterrupt
end

V(sema S)
begin
  DisableInterrupt
  s.count++
  if s.count <= 0 then
    UnblockThread(S)
  fi
  EnableInterrupt
end
```

Atomic Counting Semaphores (Multi Processor)

P(S)

begin

while (!TAS(S.flag)){}; busy waiting, flag is sema data

S.Count:= S.Count-1

if S.Count < 0 then

BlockThread(S) and S.flag := 0

else S.flag :=0

fi

end

V(S)

begin

while (!TAS(S.flag)){}; busy waiting

S.Count:= S.Count+1

if S.Count <= 0 then

UnblockThread(S)

fi

S.flag :=0

end

Next Problem Solving Exercise

